

Dheeraj Pranav

GTM AI Engineer — building auditable, evaluated agents

krovvididheeraj@gmail.com · github.com/DheerajPranav · linkedin.com/in/dheerajpranav

SUMMARY

Applied-AI engineer focused on GTM (go-to-market) systems where correctness is provable, not asserted. I build agents whose every claim is source-grounded, whose capabilities ship behind computed evaluation gates, and that report “*not measured*” rather than a fabricated number. Deep on retrieval, multi-agent orchestration, structured output, and the evals that keep them honest.

GTM AI PORTFOLIO SPRINT — SELF-DIRECTED (4 WEEKS)

Applied-AI Engineering Sprint

2026

A monorepo of four GTM AI capabilities built against one fixed, internally-consistent fictional world, so every generated claim is groundable and every capability is measured. **347 hermetic tests · \$0.00 live API spend** (every gate verified offline).

- **GTM Outbound Agent (flagship).** Built a five-agent pipeline (research → score → persona → write → critique) that returns **source-grounded** company profiles (every field carries its URL, with a deterministic check that fails fabricated provenance), scores ICP fit with a weighted mean computed in code, and drafts personalized emails via async fan-out under one shared semaphore — graded by a deliberately skeptical LLM judge. **214 hermetic tests; deterministic Account Brief.**
- **GTM Knowledge Base.** Built a hybrid-retrieval RAG assistant (BM25 + vector via Reciprocal Rank Fusion, reranker, cited answers) over a 30-doc corpus with a 35-question golden eval set. **Computed baseline: 74% hit rate@5, 61% recall@5; 84 tests.**
- **GTM Agent Evals.** Packaged a framework-agnostic, MIT-licensed LLM-judge rubric kit (ICP / Persona / Email / Critique) with deterministic gates and LangChain, external-dataset, and Streamlit integrations. **35 tests; zero framework coupling.**

Also shipped: batch mode (concurrent processing, failure isolation, checkpoint resume), a Streamlit monitoring dashboard, a Groq/Anthropic provider abstraction, and a hypothesis-driven iteration cycle. Portfolio site + repo: github.com/DheerajPranav/gtm-signal-intelligence.

TECHNICAL SKILLS

Languages	Python 3.10+, TypeScript
LLM / Agents	Anthropic Claude, Groq; tool use, multi-agent orchestration, async fan-out, prompt-injection defense
RAG	Chroma, BM25, Reciprocal Rank Fusion, rerankers, citation extraction
Evaluation	golden sets, LLM-as-judge, deterministic gates, mutation testing, Spearman/MAE metrics
Data / Backend	Pydantic v2, SQLite/Postgres, Alembic, FastAPI patterns
Frontend / Ship	Streamlit, Next.js, Tailwind; Vercel, Modal (target)
Observability	Langfuse, per-call cost/latency/token tracking

EXPERIENCE

[Add your prior roles here — company, title, dates, 2–3 impact bullets each.]

EDUCATION

[Add your degree, institution, and year.]